

Homework Assignment

HW3: Multitasking Tiny-UNIX on the SAPC

1. Introduction

The objective of this assignment is to implement a multitasking Tiny-UNIX on the SAPC. The students are asked to expand the Tiny-UNIX kernel in hw2 to multitask three programs, each of which will only make writes (to TTY0 or TTY1) and exit system calls.

In this assignment, the console output port (COM2) is the resource shared among the three programs. As in hw2, the output is interrupt-driven. This is a producer-consumer situation, with 3 producers filling the output buffer of COM2 and one consumer, the transmit interrupt handler, outputs the characters. When the output buffer is full, the processes will block in the “write” operation, giving up the CPU for other processes to use. This eliminates spin loops in the kernel.

The students will make use of the provided CPU context switching program (asmstsch.s) in writing a non-preemptive scheduler (one which never grabs the CPU away from a process which can be run) to control the running of the processes. Process switching will only occur when a process blocks (for output) or exits. After a process exits, it is a "zombie" in UNIX parlance: all that is really left of it is its exit value, waiting to be picked up by its parent. Since there is no parent process, the students will simply print out the exit values with the shutdown message of the kernel, after the completion of all three processes.

2. Discussions

The user program (here a bundle of 3 user programs) to be run on the SAPC has to be built in the same way as in hw2, downloaded and run. As in hw2, the user process is running in supervisor mode, sharing a common stack area between user and kernel for each process. There will be separate stacks for each of the three processes. The program begins with startup0.s, then startup.c. Then it calls the kernel initialization, which sets up the three processes' initial data. We will follow UNIX in reserving process 0 for the kernel itself, and require that it never block. It is ready to run at all times and if no user process has anything to do, the kernel process just loops until it exits.

A common process entry data structure, PEntry, is used to keep track of the process data:

```
typedef struct {  
    int p_savedregs[N_SAVED_REGS];    /* saved non-scratch registers */  
    ProcStatus p_status;               /* RUN, BLOCKED, or ZOMBIE */  
    WaitCode p_waitcode;              /* valid only if status=BLOCKED: TTY0_OUT, etc. */  
    int p_exitval;                    /* valid only if status=ZOMBIE */  
} PEntry;
```

In the beginning, the processes' statuses are marked as RUN. Their saved-pc entries are initialized to ustart1, ustart2 or ustart3, respectively; their saved-esp entries to three different stack areas; and their savedebp field to 0 (it controls backtrace during debugging).

Once the initialization phase is complete, the process 0 settles down to a loop calling the scheduler, trying to find and schedule a user process that can be run. But, if there are none, it keeps looping. The

Homework Assignment

loop ends when the number of ZOMBIE processes equals the number of user processes. In addition to calling the scheduler, it turns on interrupts in the loop and then off again (if interrupts are never turned on, the queue can't be drained and blocked processes will remain blocked). The scheduler should be called with interrupts off, as it is clearly critical section code. Once the main loop ends, process 0 should go into a wait loop to make sure the output queue drains and then do some finishing up, printing the exit code of each user process.

The scheduler makes use of the supplied `asmswtch.s` to perform a process switch. We'll go over the assembler code in lecture. This will in turn calls the respective C user module with entry point `main1`, `main2`, or `main3`. The syscalls in the user code cause execution of the kernel, eventually returning to the user code. However, when `ttywrite` encounters a full output buffer, it blocks the process by calling a function, `sleep(OUTPUT)`, in the scheduler (status changed from `RUN` to `BLOCKED`). The scheduler saves that process context and chooses another process to run; using the assembler `asmswtch.s` code to do the tricky stuff. When `asmswtch` returns, the system should be running the new process.

As the output drains at the interrupt (i.e., via the interrupt handler) and a spot becomes available in the output queue, a `wakeup(OUTPUT)` function is invoked which will run the processes that have been blocked on `waitcode=OUTPUT` (status changed from `BLOCKED` to `RUN`)

Finally the user code does a sys call `exit`. The kernel sets the process state to `ZOMBIE` and calls `schedule` and that finds another process to run.

3. Files

Copy the files from `hw2/soln/` and move them to your `hw3` directory. Add the files provided in my `hw3` directory to yours.

Shared between user and kernel:

<code>tsyscall.h:</code>	syscall numbers, as in <code>hw2</code>
<code>tty_public.h:</code>	TTY device numbers, as in <code>hw2</code>

Kernel files:

<code>ioconf.h, tty.h:</code>	I/O headers used only by kernel, as in <code>hw2</code>
<code>tsystm.h:</code>	syscall dispatch, kernel function prototypes as in <code>hw2</code>
<code>startup0.s, startup.c, sysentry.s:</code>	assembly & C startup, syscall handler, as in <code>hw2</code>
<code>tunix.c:</code>	kernel startup, shutdown, C trap handler, much like <code>hw2</code>
<code>sched.c</code> and <code>sched.h:</code>	scheduler code: functions <code>schedule</code> , <code>sleep</code> , <code>wakeup</code> .
<code>asmswtch.s:</code>	process switch, in assembly, provided.
<code>ioconf.c, io.c:</code>	device independent I/O system, as in <code>hw2</code>
<code>tty.c:</code>	terminal driver, from <code>hw2</code> but changed to <code>block</code> , <code>unblock</code>

User-level files:

<code>tunistd.h:</code>	syscall prototypes, as in <code>hw2</code>
<code>crt01.s:</code>	as in <code>hw2 crt0.s</code> , but calls <code>main1</code> , starts at <code>ustart1</code>
<code>crt02.s, crt03.s:</code>	calls <code>main2</code> , <code>main3</code> , starts at <code>ustart2</code> , <code>ustart3</code> .
<code>uprog1.c uprog2.c, uprog3.c:</code>	tiny-UNIX programs starting at <code>main1</code> , <code>main2</code> , <code>main3</code>
<code>ulib.s:</code>	same as <code>hw2</code>

Homework Assignment

makefile: name the executable `tunix.lnx`. Be sure to add a new rule for `sched.o` and update the rule for `tunix.o` to reflect their header inclusions. Add new rules for other new files too.

4. Suggested steps

1. Make sure your hw2 solution is fully working, or use the provided solution. Note that the provided solution has a "debug_log" service that writes notes to memory, and prints out the log when the kernel shuts down. Add this logging functionality to your own solution if you're using it for hw3. Note that you can add your own entries in the log to figure out what's happening.
2. Write trivial user programs that just `kprintf` a message, and the `crt0`'s for them that call `mainx` (where `x=1, 2, or 3`), and then do an `exit` syscall. Write a fake scheduler that just loops through the process entry table until it finds a `status = RUN` process and then calls it at `ustart1, 2 or 3`. If there aren't any processes with `status=RUN` left, bring down the system. Inside `sysexit`, call the scheduler after marking the process `ZOMBIE`. This simple system will work because each process only needs to run once. They use the same stack area one after another. Make sure your `makefile` is right: see that it rebuilds the right things after each edit.
3. Now use the supplied `asmswch` code to start each process. For each user process, you need to initialize the `PEntry`'s `saved-pc` to `ustartx` ((where `x=1, 2, or 3`), and the `saved-eflags` to allow interrupts, and the `saved-esp` to the proper stack. Also the `saved-ebx` should be 0. In your new fake scheduler, loop through the process entry table until you find a `RUN` user process (1, then 2, then 3) and then call `asmswch.s` with `old-process = process 0`, `new-process = proc 1`, to switch from `proc 0` to `proc 1`, thus running `uprog1`. Later, when `proc 1` exits, `sysexit` calls the scheduler again, which finds `proc 2`, and calls `asmswch.s` with pointers for (`proc1, proc2`) to switch from `proc 1` to `proc 2`, and that runs `uprog2`. When `uprog2` exits, `sysexit` calls the scheduler... Eventually the scheduler calls `asmswch.s` with (`proc3, proc0`) to switch back to `proc 0` when no more `RUN` processes exist. Then you can shut down the system using process 0. (Of course you can do just part of this to start, then write a more finished product.) Make sure the user processes are each using their own stack.
4. Add a `debug_log` call to the scheduler to report on each process switch, for example, "`| (2z-3)`" for a switch from process 2, a zombie, to process 3. With just `kprintf`'s in your user programs, your debug log should look like this now:

`| (0-1) | (1z-2) | (2z-3) | (3z-0)`
5. Add writes of one or two chars each to the three user programs and get this working--this output all fits in the 6-char queue, so no blocking is needed. This should work even though you haven't yet edited `tty.c`. Make sure the kernel waits for the output to drain while it's shutting down. Also make sure the user processes are running with interrupts enabled and using their appointed stack. You can check `IF` with `get_eflags()`, for example. Since the three writes return immediately, all the output will be done in the final kernel wait-for-output, and the debug log will look something like this (for 2 output chars for each process):

`| (0-1) ~ | (1z-2) ~ | (2z-3) | (3z-0) ~ ~ ~ ~ s`
(Possibly more `~s` will come earlier.)

Homework Assignment

6. Write the real scheduler after the class on this subject. Get it working in this easy user program environment. Change the code in tty.c to block and unblock at the right moments. Change to a 7-char output from one user program as soon as you want to try out blocking. When that works, try two, then all three with over-6 chars. Then try it with the provided uprog's.

The following files are the deliverables for this assignment:

README--authorship

tsystm.h--possibly new prototypes here

tunix.c--kernel initialization, shutdown, process 0 added

sched.c and sched.h--scheduler

tty.c--calls to sleep and wakeup added

crt01.s, uprog1.c

crt02.s, uprog2.c

crt03.s, uprog3.c

syms

tunix.lnx

typescript